

Stochastic Gradient Descent

Hoàng Nam Dũng

Khoa Toán - Cơ - Tin học, Đại học Khoa học Tự nhiên, Đại học Quốc gia Hà Nội

Last time: proximal gradient descent

Consider the problem

$$\min_x g(x) + h(x)$$

with g, h convex, g differentiable, and h “simple” in so much as

$$\text{prox}_t(x) = \operatorname{argmin}_z \frac{1}{2t} \|x - z\|_2^2 + h(z)$$

is computable.

Proximal gradient descent: let $x^{(0)} \in \mathbb{R}^n$, repeat:

$$x^{(k)} = \text{prox}_{t_k}(x^{(k-1)} - t_k \nabla g(x^{(k-1)})), \quad k = 1, 2, 3, \dots$$

Step sizes t_k chosen to be fixed and small, or via backtracking.

If ∇g is Lipschitz with constant L , then this has convergence rate $O(1/\varepsilon)$. Lastly we can **accelerate** this, to optimal rate $O(1/\sqrt{\varepsilon})$.

Today:

- ▶ Stochastic gradient descent
- ▶ Convergence rates
- ▶ Mini-batches
- ▶ Early stopping

Stochastic gradient descent

Consider minimizing an average of functions

$$\min_x \frac{1}{m} \sum_{i=1}^m f_i(x).$$

As $\nabla \sum_{i=1}^m f_i(x) = \sum_{i=1}^m \nabla f_i(x)$, gradient descent would repeat

$$x^{(k)} = x^{(k-1)} - t_k \cdot \frac{1}{m} \sum_{i=1}^m \nabla f_i(x^{(k-1)}), \quad k = 1, 2, 3, \dots$$

In comparison, **stochastic gradient descent** or SGD (or incremental gradient descent) repeats:

$$x^{(k)} = x^{(k-1)} - t_k \cdot \nabla f_{i_k}(x^{(k-1)}), \quad k = 1, 2, 3, \dots$$

where $i_k \in \{1, \dots, m\}$ is some chosen index at iteration k .

Stochastic gradient descent

Two rules for choosing index i_k at iteration k :

- ▶ **Randomized rule**: choose $i_k \in \{1, \dots, m\}$ uniformly at random.
- ▶ **Cyclic rule**: choose $i_k = 1, 2, \dots, m, 1, 2, \dots, m, \dots$

Randomized rule is more common in practice. For randomized rule, note that

$$\mathbb{E}[\nabla f_{i_k}(x)] = \nabla f(x),$$

so we can view SGD as using an **unbiased estimate** of the gradient at each step.

Main appeal of SGD:

- ▶ Iteration cost is independent of m (number of functions).
- ▶ Can also be a big savings in terms of memory usage.

Example: stochastic logistic regression

Given $(x_i, y_i) \in \mathbb{R}^p \times \{0, 1\}$, $i = 1, \dots, n$, recall **logistic regression**

$$\min_{\beta} f(\beta) = \frac{1}{n} \sum_{i=1}^n \underbrace{\left(-y_i x_i^T \beta + \log(1 + \exp(x_i^T \beta)) \right)}_{f_i(\beta)}.$$

Gradient computation $\nabla f(\beta) = \frac{1}{n} \sum_{i=1}^n (y_i - p_i(\beta)) x_i$ is doable when n is moderate, but **not when n is huge**.


Full gradient (also called batch) versus stochastic gradient:

- ▶ One batch update costs $O(np)$.
- ▶ One stochastic update costs $O(p)$.

Clearly, e.g., 10K stochastic steps are much more affordable.

Batch vs. stochastic gradient descent

Small example with $n = 10, p = 2$ to show the “classic picture” for batch versus stochastic methods:



Blue: batch steps, $O(np)$

Red: stochastic steps, $O(p)$

Rule of thumb for stochastic methods:

- ▶ generally thrive far from optimum
- ▶ generally struggle close to optimum

Step sizes

Standard in SGD is to use **diminishing step sizes**, e.g., $t_k = 1/k$, for $k = 1, 2, 3, \dots$

Why not fixed step sizes? Here's some intuition.

Suppose we take cyclic rule for simplicity. Set $t_k = t$ for m updates in a row, we get

$$x^{(k+m)} = x^{(k)} - t \sum_{i=1}^m \nabla f_i(x^{(k+i-1)}).$$

Meanwhile, full gradient with step size t would give

$$x^{(k+1)} = x^{(k)} - t \sum_{i=1}^m \nabla f_i(x^{(k)}).$$

The difference here: $t \sum_{i=1}^m [\nabla f_i(x^{(k+i-1)}) - \nabla f_i(x^{(k)})]$, and if we hold t constant, this difference will not generally be going to zero.

Convergence rates

Recall: for convex f , (sub)gradient descent with diminishing step sizes satisfies

$$f(x^{(k)}) - f^* = O(1/\sqrt{k}).$$

When f is differentiable with Lipschitz gradient, there holds for gradient descent with suitable fixed step sizes

$$f(x^{(k)}) - f^* = O(1/k).$$

What about SGD? For convex f , SGD with diminishing step sizes satisfies¹

$$\mathbb{E}[f(x^{(k)})] - f^* = O(1/\sqrt{k}).$$

Unfortunately this **does not improve** when we further assume f has Lipschitz gradient.

¹E.g., Nemirovski et al. (2009), “Robust stochastic optimization approach to stochastic programming”

Convergence rates

Even worse is the following discrepancy!

When f is strongly convex and has a Lipschitz gradient, gradient descent satisfies

$$f(x^{(k)}) - f^* = O(c^k)$$

where $c < 1$. But under same conditions, SGD gives us²

$$\mathbb{E}[f(x^{(k)})] - f^* = O(1/k).$$

So stochastic methods do not enjoy the **linear convergence rate** of gradient descent under strong convexity.

What can we do to improve SGD?

²E.g., Nemirovski et al. (2009), “Robust stochastic optimization approach to stochastic programming”

Mini-batches

Also common is mini-batch stochastic gradient descent, where we choose a random subset $I_k \subseteq \{1, \dots, m\}$ of size $|I_k| = b \ll m$ and repeat

$$x^{(k)} = x^{(k-1)} - t_k \cdot \frac{1}{b} \sum_{i \in I_k} \nabla f_i(x^{(k-1)}), \quad k = 1, 2, 3, \dots$$

Again, we are approximating full gradient by an unbiased estimate

$$\mathbb{E} \left[\frac{1}{b} \sum_{i \in I_k} \nabla f_i(x) \right] = \nabla f(x).$$

Using mini-batches reduces the **variance** of our gradient estimate by a factor $1/b$, but is also b times more expensive.

Batch vs. mini-batches vs. stochastic

Back to logistic regression, let's now consider a regularized version:

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{n} \sum_{i=1}^n \left(-y_i x_i^T \beta + \log(1 + e^{x_i^T \beta}) \right) + \frac{\lambda}{2} \|\beta\|_2^2.$$

Write the criterion as

$$f(\beta) = \frac{1}{n} \sum_{i=1}^n f_i(\beta), \quad f_i(\beta) = -y_i x_i^T \beta + \log(1 + e^{x_i^T \beta}) + \frac{\lambda}{2} \|\beta\|_2^2.$$


Full gradient computation is $\nabla f(\beta) = \frac{1}{n} \sum_{i=1}^n (y_i - p_i(\beta)) x_i + \lambda \beta$.

Comparison between methods:

- ▶ One batch update costs $O(np)$.
- ▶ One mini-batch update costs $O(bp)$.
- ▶ One stochastic update costs $O(p)$.


Batch vs. mini-batches vs. stochastic

Example with $n = 10,000$, $p = 20$, all methods use fixed step sizes



Batch vs. mini-batches vs. stochastic


What's happening? Now let's parametrize by flops³



³flops: floating point operations per second

Batch vs. mini-batches vs. stochastic

Finally, looking at suboptimality gap (on log scale)



End of the story?

Short story:

- ▶ SGD can be **super effective** in terms of iteration cost, memory.
- ▶ But SGD is **slow to converge**, can't adapt to strong convexity.
- ▶ And mini-batches seem to be a wash in terms of flops (though they can still be useful in practice).

End of the story?

Short story:

- ▶ SGD can be **super effective** in terms of iteration cost, memory.
- ▶ But SGD is **slow to converge**, can't adapt to strong convexity.
- ▶ And mini-batches seem to be a wash in terms of flops (though they can still be useful in practice).

Is this the end of the story for SGD?

End of the story?

Short story:

- ▶ SGD can be **super effective** in terms of iteration cost, memory.
- ▶ But SGD is **slow to converge**, can't adapt to strong convexity.
- ▶ And mini-batches seem to be a wash in terms of flops (though they can still be useful in practice).

Is this the end of the story for SGD?

For a while, the answer was believed to be yes. Slow convergence for strongly convex functions was believed inevitable, as Nemirovski and others established matching **lower bounds** ... but this was for a more general stochastic problem, where $f(x) = \int F(x, \zeta) dP(\zeta)$.

New wave of “variance reduction” work shows we can modify SGD to converge much faster for finite sums (more later?).

SGD in large-scale ML

SGD has really taken off in large-scale machine learning

- ▶ In many ML problems we don't care about optimizing to high accuracy, it doesn't pay off in terms of statistical performance.
- ▶ Thus (in contrast to what classic theory says) **fixed step sizes** are commonly used in ML applications.
- ▶ One trick is to experiment with step sizes using small fraction of training before running SGD on full data set ... many other heuristics are common⁴.
- ▶ Many variants provide better practical stability, convergence: momentum, acceleration, averaging, coordinate-adapted step sizes, variance reduction ...
- ▶ See AdaGrad, Adam, AdaMax, SVRG, SAG, SAGA ... (more later?).

⁴E.g., Bottou (2012), "Stochastic gradient descent tricks"

Early stopping

Suppose p is large and we wanted to fit (say) a logistic regression model to data $(x_i, y_i) \in \mathbb{R}^p \times \{0, 1\}$, $i = 1, \dots, n$.

We could solve (say) ℓ_2 regularized logistic regression

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{n} \sum_{i=1}^n \left(-y_i x_i^T \beta + \log(1 + e^{x_i^T \beta}) \right) \text{ subject to } \|\beta\|_2 \leq t$$

We could also run gradient descent on the unregularized problem

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{n} \sum_{i=1}^n \left(-y_i x_i^T \beta + \log(1 + e^{x_i^T \beta}) \right)$$

and **stop early**, i.e., terminate gradient descent well-short of the global minimum.

Early stopping

Consider the following, for a very small constant step size ε :

- ▶ Start at $\beta^{(0)} = 0$, solution to regularized problem at $t = 0$.
- ▶ Perform gradient descent on unregularized criterion

$$\beta^{(k)} = \beta^{(k-1)} - \varepsilon \cdot \frac{1}{n} \sum_{i=1}^n (y_i - p_i(\beta^{(k-1)})) x_i, \quad k = 1, 2, 3, \dots$$

(we could equally well consider SGD).


- ▶ Treat $\beta^{(k)}$ as an approximate solution to regularized problem with $t = \|\beta^{(k)}\|_2$.

This is called **early stopping** for gradient descent. Why would we ever do this? It's both more convenient and potentially much more efficient than using explicit regularization.

An intriguing connection

When we solve the ℓ_2 regularized logistic problem for varying t , solution path looks quite similar to gradient descent path!

Example with $p = 8$, solution and grad descent paths side by side






Lots left to explore

- ▶ Connection holds beyond logistic regression, for arbitrary loss.
- ▶ In general, the grad descent path will not coincide with the ℓ_2 regularized path (as $\varepsilon \rightarrow 0$). Though in practice, it seems to give competitive statistical performance.
- ▶ Can extend early stopping idea to mimic a generic regularizer (beyond ℓ_2)⁵.
- ▶ There is a lot of literature on early stopping, but it's still not as well-understood as it should be.
- ▶ Early stopping is just one instance of **implicit** or **algorithmic regularization**, many others are effective in large-scale ML, they all should be better understood

⁵Tibshirani (2015), "A general framework for fast stagewise algorithms"

References and further reading

-  D. Bertsekas (2010), *Incremental gradient, subgradient, and proximal methods for convex optimization: a survey*
-  A. Nemirovski and A. Juditsky and G. Lan and A. Shapiro (2009), *Robust stochastic optimization approach to stochastic programming*
-  R. Tibshirani (2015), *A general framework for fast stagewise algorithms*